

---

# Classifying Handwritten Digits with a Convolutional Deep Neural Network

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1        Neural networks excel at solving classification problems. A simple implementation  
2        of a deep neural network with convolutional and pooling layers saw success at  
3        classifying members of the MNIST dataset: labeled images of handwritten digits.

## 4    1    Introduction

5        In order to classify the digits in the MNIST dataset, a deep convolutional neural network was used,  
6        but what exactly does that mean? The network contained convolutional, pooling, and fully connected  
7        layers, each serving either to increase efficiency of training or accuracy of classification. Training was  
8        performed on randomly selected batches of 100 digits. 5000 of these batches were used for training.  
9        This process was repeated twice more in order to achieve a clearer idea of the networks performance.

## 10   2    The MNIST Dataset

11       The MNIST dataset consists of 70,000 images of handwritten digits, and labels defining which digit  
12       they represent. There are 60,000 in the training set, and 10,000 in the testing set. Each image has a  
13       height and width of 28 pixels on a single black/white color channel.

## 14   3    Network Structure

15       The network consisted of of two convolutional layers, two pooling layers, and one fully connected  
16       layer.

### 17   3.1   Convolutional Layers

18       Each convolutional layer applied a number of convolution filters to 5x5 subregions of the image. This  
19       forced the neural network to consider regions of the image in the same way a human might from its  
20       start, instead of allowing it to learn that groups of nearby pixels have significance over many training  
21       steps.

22       The first convolutional layer applied 16 filters, which multiplied the size of the data by 16. The  
23       second applied 32 filters, which similarly increased the size by 32.

24       The second convolutional layer was acting not on pixels, but on the data generated by the first  
25       convolutional layer. Essentially, it drew inferences about the input not from relationships between  
26       groups of pixels, but from relationships between groups of those groups.

## 27 3.2 Pooling Layers

28 Between the two convolutional layers, the amount of data being processed grows massively, which  
29 greatly increases the amount of processing time required for training. This was remedied in part by  
30 the inclusion of pooling layers after each convolutional layer, which reduced the resolution of the data  
31 by a factor of four. This comes at a cost to quality of classification, but with such a great decrease in  
32 processing time, its inclusion is well worth it.

33 The specific method of pooling used was max pooling over 2x2 regions with a step size of 2. This  
34 means that the data entering the layer was broken into discrete 2x2 regions, and each was reduced to  
35 its single highest value. Another method, such as taking the average of the values within each region,  
36 or taking the most extreme value could have been used. This process can be thought of in much the  
37 same way one can think of reducing the resolution on any image. Quality is indisputably lost, but  
38 without massive reduction in resolution, the image is still much the same.

## 39 3.3 Dense Layer

40 Finally, after all convolution and pooling is complete, the data passes through one dense, or fully  
41 connected layer, which simply draws patterns from the data passed in, which is finally reduced to  
42 ten values, one for each possible label. The greatest output value acts as the network's prediction.  
43 The specific dense layer used here contained 512 nodes.

44 This layer had a dropout rate of 40%, which means that 40% of values, randomly selected, are simply  
45 not passed forwards. This encourages a robustness in the network that should allow it to succeed at  
46 classification even missing 40% of data, and should help to avoid total reliance on any one feature.

## 47 4 Training

48 In order to train the network, a batch of randomly selected training samples is selected and run  
49 through the network. Outputs are generated by the network for each member of the batch, and loss  
50 is calculated. Loss is simply the difference between the desired output and the actual output of the  
51 network. In this case, the desired output is a zero for each label except for the correct one, which  
52 should have a value of one.

53 This loss is used to perform gradient descent on the network, by which weight values are changed  
54 from the output layer working back towards the beginning, such that, should the same input arrive  
55 again, the actual output will be closer to the desired output. The learning rate for this process was  
56 selected as 0.001. This value was selected arbitrarily, and could very well be changed for better  
57 results. Empirical testing is required to find the best learning rate.

58 This process was repeated 5000 times, and then the network was reset, and learning was done again  
59 from scratch twice more.

## 60 5 Results

61 Training steps took on average 27.3 seconds each. The average accuracy after training was 93.55%.  
62 Compared to some attempts at solving the same problem with neural networks, this was a bit of a bust.  
63 As of 2013, some attempts were seeing accuracy as high as 99.79%. Very impressive. With much  
64 more training time and number fiddling based on empirical testing, the implimentation described here  
65 could possibly have reached similar quality.